

Averages, Expectation, Aggregation

Understanding Political Numbers

Feb 18, 2019

Learning ggplot

How to draw an owl

1.



1. Draw some circles

2.



2. Draw the rest of the fucking owl

(Spooky voice) *Statistiiiiiiiics*

This week: "the signal and the noise"

- Today: Means
- Wednesday: Variance

In section: major tidyverse functions

Questions about exercise 1?

Major functions in the tidyverse

a.k.a. "verbs." They modify and return *data frames*

Function	Operation
<code>arrange()</code>	Sort data frame along variable(s)
<code>select()</code>	Choose variables (columns) from a data frame
<code>filter()</code>	Choose cases (rows) from a data frame
<code>mutate()</code>	Create or modify variables
<code>count()</code>	Tabulate variable(s) in a data frame
<code>summarize()</code>	Calculate summary statistics from a data frame
<code>group_by()</code>	Implicitly partition a data frame along variable(s)

Arrange

```
# Load tidyverse and gapminder data
library("tidyverse")
library("gapminder")

# With tidyverse verbs, the first argument is the data frame
# Sort by year and then by continent.
arrange(gapminder, year, continent)
```

```
## # A tibble: 1,704 x 6
```

```
##   country      continent  year lifeExp    pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>   <int>    <dbl>
## 1 Algeria      Africa    1952   43.1  9279525   2449.
## 2 Angola        Africa    1952   30.0  4232095   3521.
## 3 Benin         Africa    1952   38.2  1738315   1063.
## 4 Botswana      Africa    1952   47.6   442308    851.
## 5 Burkina Faso   Africa    1952   32.0  4469979    543.
## 6 Burundi       Africa    1952   39.0  2445618    339.
## 7 Cameroon      Africa    1952   38.5  5009067   1173.
## 8 Central African Republic Africa    1952   35.5  1291695   1071.
## 9 Chad          Africa    1952   38.1  2682462   1179.
## 10 Comoros       Africa    1952   40.7   153936   1103.
## # ... with 1,694 more rows
```

Select variables

```
# Which variables do I want to keep?  
# Again, first arg is the data frame name  
# (notice lack of $)  
select(gapminder, country, year, gdpPercap)
```

```
## # A tibble: 1,704 x 3  
##   country      year gdpPercap  
##   <fct>      <int>     <dbl>  
## 1 Afghanistan  1952      779.  
## 2 Afghanistan  1957      821.  
## 3 Afghanistan  1962      853.  
## 4 Afghanistan  1967      836.  
## 5 Afghanistan  1972      740.  
## 6 Afghanistan  1977      786.  
## 7 Afghanistan  1982      978.  
## 8 Afghanistan  1987      852.  
## 9 Afghanistan  1992      649.  
## 10 Afghanistan 1997      635.  
## # ... with 1,694 more rows
```

Filter observations

```
# Which cases (rows) do I want to keep?  
# filter(dataset, logical test)  
# keep rows where test result is TRUE  
filter(gapminder, country == "United States")
```

Logical operators:

- `==` means "is equal to"
- `!=` means "not equal to"
- `>` and `<` mean "greater/less than"
- `>=` and `<=` are "greater than/less than or equal to"

Combine logical tests with `&` (and) or `|` (or)

- `filter(gapminder, country == "United States" & year > 2000)`

```
## # A tibble: 12 x 6  
##   country      continent  year lifeExp      pop gdpPercap  
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>  
## 1 United States Americas   1952    68.4 157553000 13990.  
## 2 United States Americas   1957    69.5 171984000 14847.  
## 3 United States Americas   1962    70.2 186538000 16173.  
## 4 United States Americas   1967    70.8 198712000 19530.  
## 5 United States Americas   1972    71.3 209896000 21806.  
## 6 United States Americas   1977    73.4 220239000 24073.  
## 7 United States Americas   1982    74.6 232187835 25010.  
## 8 United States Americas   1987    75.0 242803533 29884.  
## 9 United States Americas   1992    76.1 256894189 32004.  
## 10 United States Americas   1997    76.8 272911760 35767.  
## 11 United States Americas   2002    77.3 287675526 39097.  
## 12 United States Americas   2007    78.2 301139947 42952.
```

Create variables with "mutate"

```
# mutate(dataframe, new_variable = (whatever you want))  
mutate(gapminder,  
       gdp = gdpPercap * pop)
```

```
## # A tibble: 1,704 x 7
```

```
##   country      continent  year lifeExp      pop gdpPercap      gdp  
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>    <dbl>  
## 1 Afghanistan Asia      1952   28.8  8425333   779.  6567086330.  
## 2 Afghanistan Asia      1957   30.3  9240934   821.  7585448670.  
## 3 Afghanistan Asia      1962   32.0 10267083   853.  8758855797.  
## 4 Afghanistan Asia      1967   34.0 11537966   836.  9648014150.  
## 5 Afghanistan Asia      1972   36.1 13079460   740.  9678553274.  
## 6 Afghanistan Asia      1977   38.4 14880372   786. 11697659231.  
## 7 Afghanistan Asia      1982   39.9 12881816   978. 12598563401.  
## 8 Afghanistan Asia      1987   40.8 13867957   852. 11820990309.  
## 9 Afghanistan Asia      1992   41.7 16317921   649. 10595901589.  
## 10 Afghanistan Asia      1997   41.8 22227415   635. 14121995875.  
## # ... with 1,694 more rows
```


Count (or tabulate)

```
# tabulate variable(s) with count().  
# Again... result is a DATA FRAME  
count(gapminder, continent, year)
```

```
## # A tibble: 60 x 3  
##   continent year      n  
##   <fct>      <int> <int>  
## 1 Africa    1952     52  
## 2 Africa    1957     52  
## 3 Africa    1962     52  
## 4 Africa    1967     52  
## 5 Africa    1972     52  
## 6 Africa    1977     52  
## 7 Africa    1982     52  
## 8 Africa    1987     52  
## 9 Africa    1992     52  
## 10 Africa   1997     52  
## # ... with 50 more rows
```

Summarize variables

```
# New data frame of summary calculations  
# Use na.rm = TRUE to skip missing values when calculating summary stats  
summarize(gapminder,  
          mean_lifeexp = mean(lifeExp),  
          min_lifeexp = min(lifeExp),  
          max_lifeexp = max(lifeExp, na.rm = TRUE))
```

```
## # A tibble: 1 x 3  
##   mean_lifeexp min_lifeexp max_lifeexp  
##         <dbl>         <dbl>         <dbl>  
## 1         59.5          23.6          82.6
```

Group data by variables

```
# partition data into groups. Pretty benign when used alone  
group_by(gapminder, continent)
```

```
## # A tibble: 1,704 x 6  
## # Groups:   continent [5]  
##   country    continent  year lifeExp      pop gdpPercap  
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>  
## 1 Afghanistan Asia      1952   28.8  8425333    779.  
## 2 Afghanistan Asia      1957   30.3  9240934    821.  
## 3 Afghanistan Asia      1962   32.0 10267083    853.  
## 4 Afghanistan Asia      1967   34.0 11537966    836.  
## 5 Afghanistan Asia      1972   36.1 13079460    740.  
## 6 Afghanistan Asia      1977   38.4 14880372    786.  
## 7 Afghanistan Asia      1982   39.9 12881816    978.  
## 8 Afghanistan Asia      1987   40.8 13867957    852.  
## 9 Afghanistan Asia      1992   41.7 16317921    649.  
## 10 Afghanistan Asia      1997   41.8 22227415    635.  
## # ... with 1,694 more rows
```

Group and summarize

```
# the `<-` scans the next line
gap_by_continent <-
  group_by(gapminder, continent)

summarize(gap_by_continent,
  mean_life = mean(lifeExp))
```

```
## # A tibble: 5 x 2
##   continent mean_life
##   <fct>      <dbl>
## 1 Africa      48.9
## 2 Americas    64.7
## 3 Asia        60.1
## 4 Europe      71.9
## 5 Oceania     74.3
```

```
# Because result of group_by() is a data frame,
# you could pass result directly to summarize
summarize(group_by(gapminder, continent),
  mean_life = mean(lifeExp))
```

```
## # A tibble: 5 x 2
##   continent mean_life
##   <fct>      <dbl>
## 1 Africa      48.9
## 2 Americas    64.7
## 3 Asia        60.1
## 4 Europe      71.9
## 5 Oceania     74.3
```

Averages

Question: do women vote for Democrats more than men do?


Break the question down:

1. What's the *average* rate of Democratic voting among women?
2. Among men?
3. How different are they?

Question: do women vote for Democrats more than men do?

Break the question down:

1. What's the *average* rate of Democratic voting among women?
2. Among men?
3. How different are they?



Question: is voter turnout higher among older voters?

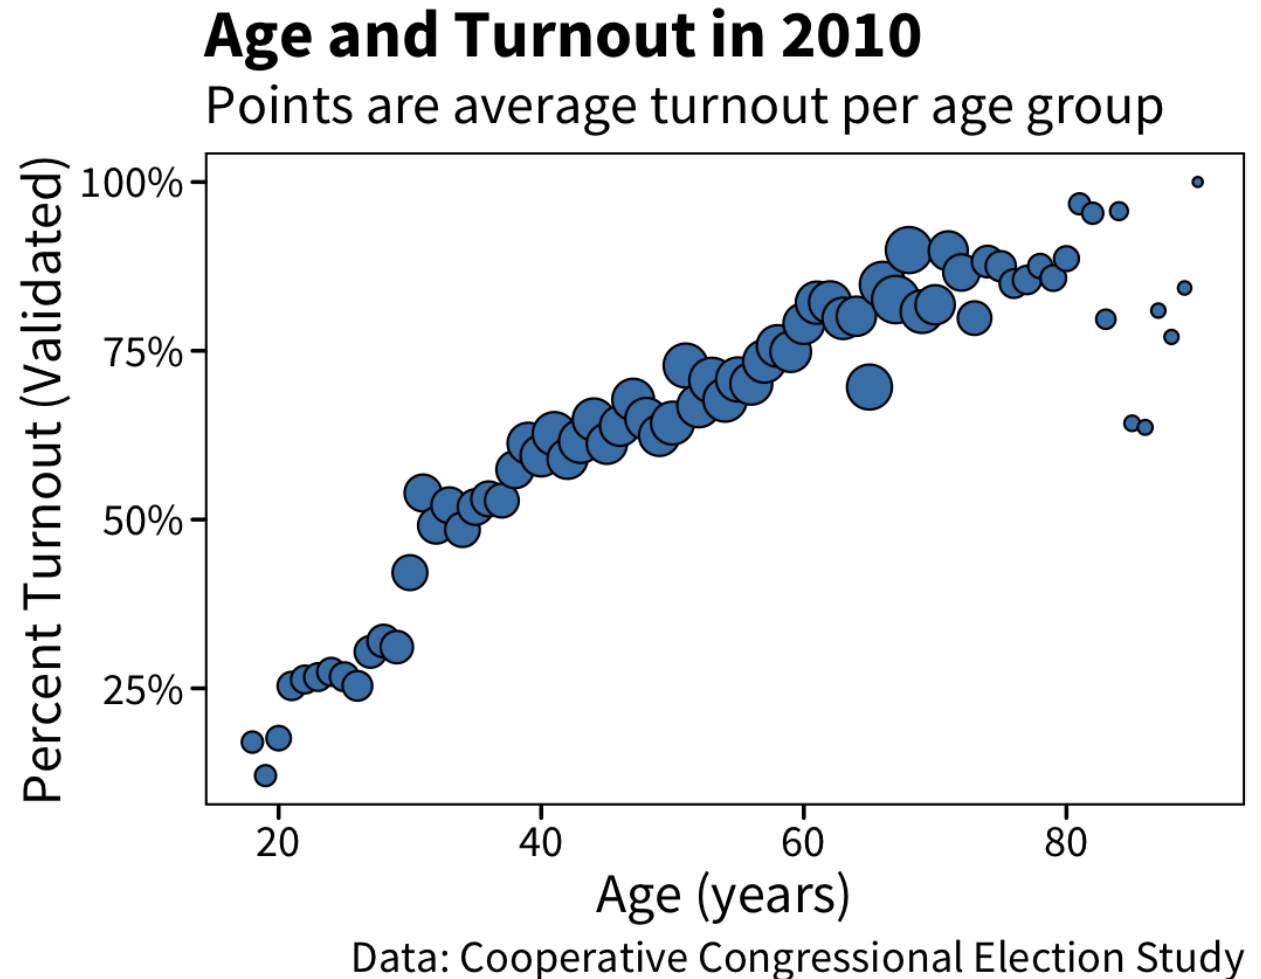
As a comparison of averages:

- Average turnout among older voters
- Among younger voters
- with a twist: age is continuous(ish)

Question: is voter turnout higher among older voters?

As a comparison of averages:

- Average turnout among older voters
- Among younger voters
- with a twist: age is continuous(ish)



Averages are useful because they tell us about
the *typical* behavior in the data

Practically & Ethically: individuals \neq averages

Averaging (the math)

$$x = [6 \quad 15 \quad 8 \quad 16 \quad 17].$$

The average of x , we call \bar{x} .

Averaging (the math)

$$x = [6 \quad 15 \quad 8 \quad 16 \quad 17].$$

The average of x , we call \bar{x} .

$$\bar{x} = \frac{\sum_{i=1}^n x_i}{n}$$

$$\bar{x} = \frac{1}{n} \sum_i^n x_i$$

Averaging (the math)

$$x = [6 \quad 15 \quad 8 \quad 16 \quad 17].$$

The average of x , we call \bar{x} .

$$\bar{x} = \frac{\sum_{i=1}^n x_i}{n}$$

$$\bar{x} = \frac{1}{n} \sum_i^n x_i$$

```
mean(x)
```

```
## [1] 12.4
```

```
sum(x) / length(x)
```

```
## [1] 12.4
```

Strategies for averaging different data types

Strategies for averaging different data types

Quantitative (interval and ratio) data

```
summarize(gapminder,  
          avg_lifeexp = mean(lifeExp),  
          avg_gdpPercap = mean(gdpPercap))
```

```
## # A tibble: 1 x 2  
##   avg_lifeexp avg_gdpPercap  
##       <dbl>       <dbl>  
## 1      59.5      7215.
```

Strategies for averaging different data types

Quantitative (interval and ratio) data

```
summarize(gapminder,  
          avg_lifeexp = mean(lifeExp),  
          avg_gdpPercap = mean(gdpPercap))
```

```
## # A tibble: 1 x 2  
##   avg_lifeexp avg_gdpPercap  
##       <dbl>       <dbl>  
## 1      59.5       7215.
```

Categorical (nominal and ordinal) data

```
# Proportion of data in continents.  
summarize(gapminder,  
          pr_afr = mean(continent == "Africa"),  
          pr_euro = mean(continent == "Europe"))
```

```
## # A tibble: 1 x 2  
##   pr_afr pr_euro  
##   <dbl> <dbl>  
## 1  0.366  0.211
```

If we have a vector of 1s and 0s ("successes" and "failures"), the mean is equal to the proportion of 1s (successes)

Why we like averages: noise canceling out

We flip a coin.

```
# make a coin vector  
coin <- c("Heads", "Tails")  
  
# "flip" the coin  
sample(coin, 1)
```

```
## [1] "Heads"
```

Why we like averages: noise canceling out

We flip a coin.

```
# make a coin vector  
coin <- c("Heads", "Tails")  
  
# "flip" the coin  
sample(coin, 1)
```

```
## [1] "Heads"
```

We flip it 5 times.

```
# 'replace' means we put the coin back each time  
flips <- sample(coin, 5, replace = TRUE)  
  
# what's the proportion of heads?  
mean(flips == "Heads")
```

```
## [1] 0.2
```

Why we like averages: noise canceling out

Flip 100 times. After each flip, find proportion of heads *up to that point*

Eventually this "running average" should approach what number?

```
## # A tibble: 100 x 3
##   trial flip  running_mean
##   <int> <chr>      <dbl>
## 1      1 1 Heads          1
## 2      2 2 Tails         0.5
## 3      3 3 Heads        0.667
## 4      4 4 Tails         0.5
## 5      5 5 Tails         0.4
## 6      6 6 Heads         0.5
## 7      7 7 Tails        0.429
## 8      8 8 Tails        0.375
## 9      9 9 Tails        0.333
## 10     10 10 Heads        0.4
## # ... with 90 more rows
```

Why we like averages: noise canceling out


Flip 100 times. After each flip, find proportion of heads *up to that point*

Eventually this "running average" should approach what number?

```
## # A tibble: 100 x 3
##   trial flip  running_mean
##   <int> <chr>      <dbl>
## 1     1  Heads         1
## 2     2  Tails         0.5
## 3     3  Heads         0.667
## 4     4  Tails         0.5
## 5     5  Tails         0.4
## 6     6  Heads         0.5
## 7     7  Tails         0.429
## 8     8  Tails         0.375
## 9     9  Tails         0.333
## 10    10 Heads         0.4
## # ... with 90 more rows
```

Long-run average of coin flips

... after n trials



Expectation

Expectation

The true / theoretical / long-run average

Expectation

The true / theoretical / long-run average

Example: more coins

Suppose that the variable \mathbf{X} contains an *arbitrary number* of coin flips (1 = "Heads", 0 = "Tails").

As the number of trials approaches $+\infty$, the mean of \mathbf{X} approaches what value?

Expectation

The true / theoretical / long-run average

Example: more coins

Suppose that the variable \mathbf{X} contains an *arbitrary number* of coin flips (1 = "Heads", 0 = "Tails").

As the number of trials approaches $+\infty$, the mean of \mathbf{X} approaches what value?

$$E[\mathbf{X}] = 0.5$$

Probabilities are one type of *expectation*, but not the only kind

Another toy example: rolling a die

Roll a six-sided die. What's the expected value?

Another toy example: rolling a die

Roll a six-sided die. What's the expected value? It's 3.5, why?

Another toy example: rolling a die

Roll a six-sided die. What's the expected value? It's 3.5, why?

It's the "theoretical average."

Another toy example: rolling a die

Roll a six-sided die. What's the expected value? It's 3.5, why?

It's the "theoretical average."

$$\begin{aligned} E[X] &= \sum_{k=1}^K x_k p_k \\ &= x_1 p_1 + x_2 p_2 + \dots + x_K p_K \end{aligned}$$

- x_k represents a possible outcome
- p_k is the probability of outcome k
- K is the total number of possibilities

Another toy example: rolling a die

Roll a six-sided die. What's the expected value? It's **3.5**, why?

It's the "theoretical average."

$$\begin{aligned} E[X] &= \sum_{k=1}^K x_k p_k \\ &= x_1 p_1 + x_2 p_2 + \dots + x_K p_K \end{aligned}$$

- x_k represents a possible outcome
- p_k is the probability of outcome k
- K is the total number of possibilities

$$E[\text{Die}] = \left(1 \times \frac{1}{6}\right) + \left(2 \times \frac{1}{6}\right) + \dots + \left(6 \times \frac{1}{6}\right) = 3.5$$

Expectation is a **weighted average of all possible outcomes** (each outcome weighted by its probability of occurrence)

Why does this matter?

The *theoretical average* influences the data we collect, but our data don't *perfectly reflect* the true expectation

Why does this matter?

The *theoretical average* influences the data we collect, but our data don't *perfectly reflect* the true expectation

- Candidate A and B are running for Senate
- TRUE support for Candidate A is 54%
- "the population parameter" ($\mu = 0.54$)

Why does this matter?

The *theoretical average* influences the data we collect, but our data don't *perfectly reflect* the true expectation

- Candidate A and B are running for Senate
- TRUE support for Candidate A is 54%
- "the population parameter" ($\mu = 0.54$)

We take a survey of 500 voters

```
# Voters support A or B randomly with given probabilities.  
# repeat 500x with replacement  
voters <- sample(c("A", "B"), prob = c(0.54, 1 - 0.54),  
                size = 500, replace = TRUE)  
  
# what proportion of the sample is voting for A?  
mean(voters == "A")
```

```
## [1] 0.56
```


Why does this matter?

The *theoretical average* influences the data we collect, but our data don't *perfectly reflect* the true expectation

- Candidate A and B are running for Senate
- TRUE support for Candidate A is 54%
- "the population parameter" ($\mu = 0.54$)

We take a survey of 500 voters

```
# Voters support A or B randomly with given probabilities.  
# repeat 500x with replacement  
voters <- sample(c("A", "B"), prob = c(0.54, 1 - 0.54),  
                size = 500, replace = TRUE)  
  
# what proportion of the sample is voting for A?  
mean(voters == "A")
```

```
## [1] 0.56
```

Expected value (a.k.a. "population mean") is $\mu = 0.54$

Estimated value (a.k.a. "sample mean") is $\bar{x} = 0.56$

Taking a sample

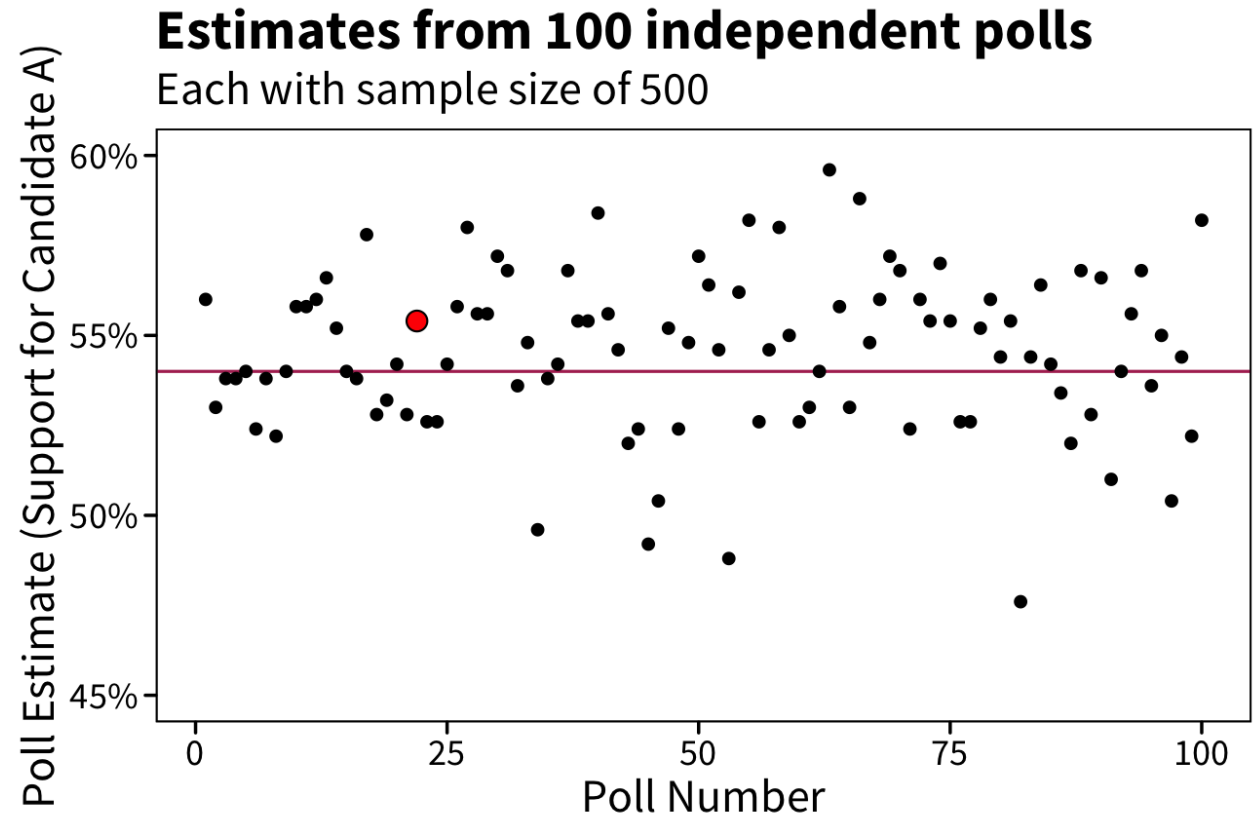
We will never know the *true mean* (μ) with certainty.

But we can take samples of data and calculate the *sample mean* (\bar{x}) within the sample.

Taking a sample

We will never know the *true mean* (μ) with certainty.

But we can take samples of data and calculate the *sample mean* (\bar{x}) within the sample.



The whole point of statistics
is to figure out how confident we can be about the *real*
truth

given that we only can observe our imperfect sample of data

Aggregation

Aggregation

We have data at some level of analysis, and we want to summarize it at a higher level of analysis

Aggregation

We have data at some level of analysis, and we want to summarize it at a higher level of analysis

Life expectancy *aggregated* by year

```
## # A tibble: 12 x 2
##   year lifeExp_mean
##   <int>         <dbl>
## 1  1952          49.1
## 2  1957          51.5
## 3  1962          53.6
## 4  1967          55.7
## 5  1972          57.6
## 6  1977          59.6
## 7  1982          61.5
## 8  1987          63.2
## 9  1992          64.2
## 10 1997          65.0
## 11 2002          65.7
## 12 2007          67.0
```

Aggregation

We have data at some level of analysis, and we want to summarize it at a higher level of analysis


Life expectancy *aggregated* by year

```
## # A tibble: 12 x 2
##   year lifeExp_mean
##   <int>         <dbl>
## 1  1952          49.1
## 2  1957          51.5
## 3  1962          53.6
## 4  1967          55.7
## 5  1972          57.6
## 6  1977          59.6
## 7  1982          61.5
## 8  1987          63.2
## 9  1992          64.2
## 10 1997          65.0
## 11 2002          65.7
## 12 2007          67.0
```

Life expectancy *aggregated* by continent-year

```
## # A tibble: 60 x 3
## # Groups:   continent [5]
##   continent year lifeExp_mean
##   <fct>     <int>         <dbl>
## 1 Africa    1952          39.1
## 2 Africa    1957          41.3
## 3 Africa    1962          43.3
## 4 Africa    1967          45.3
## 5 Africa    1972          47.5
## 6 Africa    1977          49.6
## 7 Africa    1982          51.6
## 8 Africa    1987          53.3
## 9 Africa    1992          53.6
## 10 Africa   1997          53.6
## # ... with 50 more rows
```


"Conditional" averages and "conditional" probabilities



THE ANNUAL DEATH RATE AMONG PEOPLE
WHO KNOW THAT STATISTIC IS ONE IN SIX.

Ecological Fallacy:

Assuming that group-level patterns apply to individuals within the group


Obama-Trump voters turn back to Democrats

Senate Democrats did especially well where Donald Trump had gained the most ground

Swing and a prayer

United States, by county

Democratic share of votes for senator (2018) minus
Hillary Clinton's vote share (2016), percentage points



Source: Edison Research

The Economist

See ya

Wednesday is (spooky voice) randomnessssssss